

Advanced topics of AI

Large Language Models





[Adapted from slides by Dan Klein and Pieter Abbeel at UC Berkeley]

Advanced topics of AI

- Large language models
- Advanced reinforcement learning
- Advanced deep learning
- Responsible AI


Today's AI



Large Language Models

- Feature engineering
 - Text tokenization
 - Word embeddings
- Deep neural networks
 - Autoregressive models
 - Self-attention mechanisms
 - Transformer architecture
- Multi-class classification
- Supervised learning
 - Self-supervised learning
 - Instruction tuning
- Reinforcement learning
 - ... from human feedback (RLHF)
- Policy search
 - Policy gradient methods
- Beam search

Deep Neural Networks



- Input: some text

- “The dog chased the”

- Output: more text

-

... “ball”

- Implementation:

- Linear algebra
- How??

Text Tokenization

GPT-3.5 & GPT-4

GPT-3 (Legacy)

Many words map to one token, but some don't: indivisible.

Unicode characters like emojis may be split into many tokens containing the underlying bytes: 🖐

Sequences of characters commonly found next to each other may be grouped together: 1234567890

Clear

Show example

Tokens
57

Characters
252

Text Tokenization

GPT-3.5 & GPT-4

GPT-3 (Legacy)

Many words map to one token, but some don't: indivisible.

Unicode characters like emojis may be split into many tokens containing the underlying bytes: 🍌🍌🍌🍌🍌

Sequences of characters commonly found next to each other may be grouped together: 1234567890

Text

Token IDs

Tokens

57

Characters

252

Text Tokenization

GPT-3.5 & GPT-4

GPT-3 (Legacy)

```
[8607, 4339, 2472, 311, 832, 4037, 11, 719, 1063, 1541, 956, 25, 3687,
23936, 382, 35020, 5885, 1093, 100166, 1253, 387, 6859, 1139, 1690,
11460, 8649, 279, 16940, 5943, 25, 11410, 97, 248, 9468, 237, 122, 271,
1542, 45045, 315, 5885, 17037, 1766, 1828, 311, 1855, 1023, 1253, 387,
41141, 3871, 25, 220, 4513, 10961, 16474, 15]
```

Text

Token IDs

Tokens

57


Characters

252

Word Embeddings


- Input: some text

- “The”
- “ dog”
- “ chased”
- “ the”




one-hot

[791]
[5679]
[62920]
[279]





- Output: more text

- “ ball”




[5041]




What do word embeddings look like?

- Words cluster by similarity:




What do word embeddings look like?

- Features learned in language models:



What do word embeddings look like?

- Signs of sensible algebra in embedding space:



[Efficient estimation of word representations in vector space, Mikolov et al, 2013]

Aside: interactive explainer of modern language models

ig.ft.com/generative-ai

Artificial Intelligence

Generative AI exists because of the transformer

This is how it works

By Visual Storytelling Team and Madhumita Murgia in London SEPTEMBER 11 2023

Large Language Models

- ~~Feature engineering~~

- ~~Text tokenization~~

- ~~Word embeddings~~

- Deep neural networks

- Autoregressive models

- Self-attention mechanisms

- Transformer architectures

- Multi-class classification

- Supervised learning

- Self-supervised learning

- Instruction tuning

- Reinforcement learning

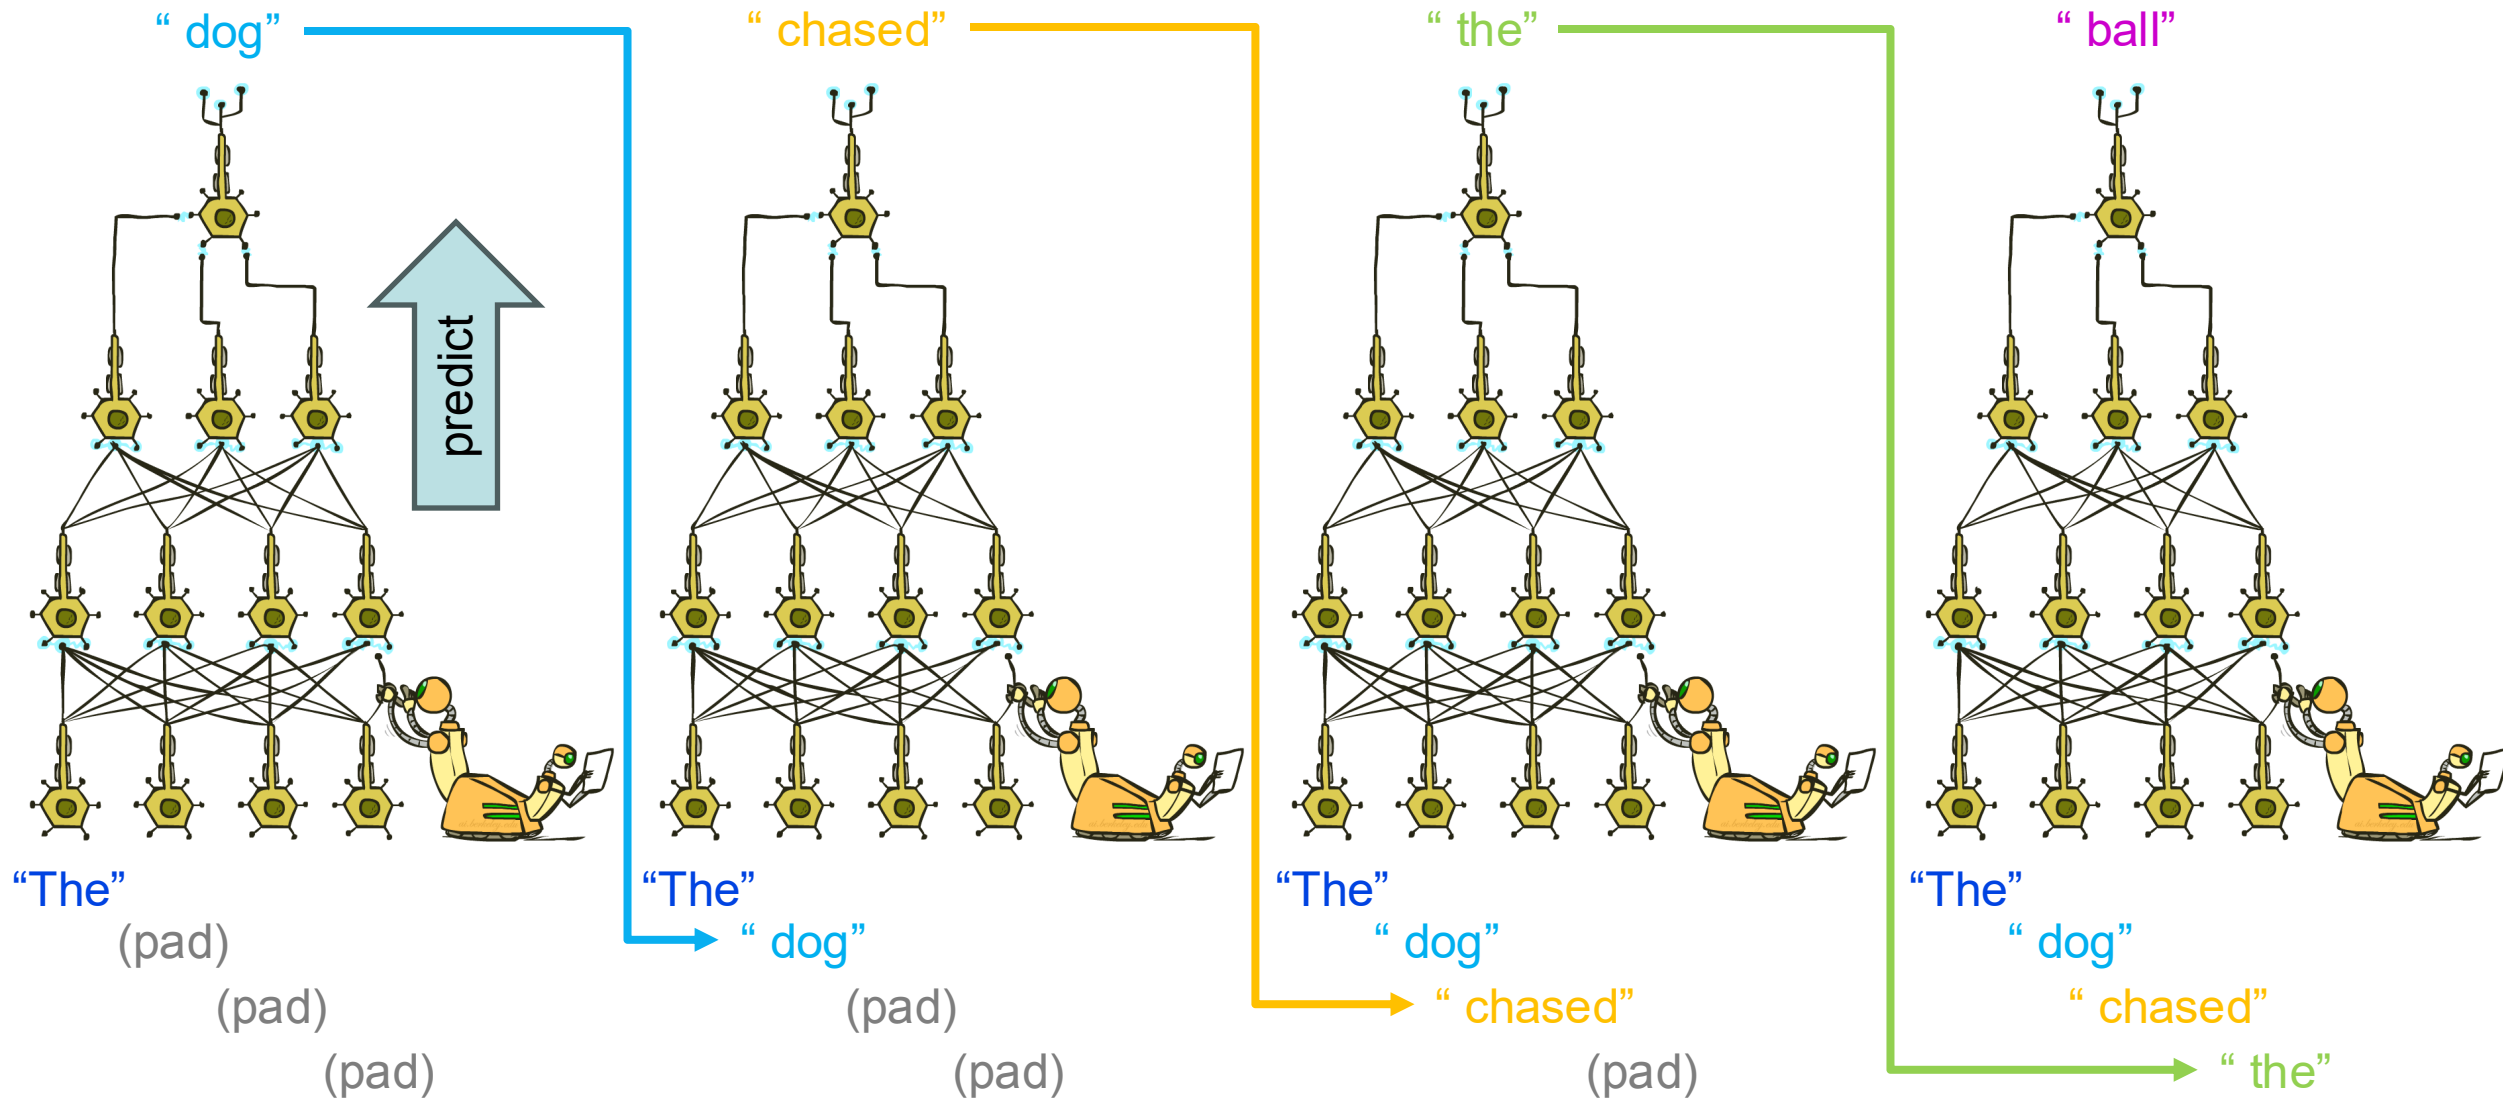
- ... from human feedback (RLHF)

- Policy search

- Policy gradient methods


- Beam search

Autoregressive Models




Autoregressive Models

- Predict output one piece at a time (e.g. word, token, pixel, etc.)
- Concatenate: input + output
- Feed result back in as new input
- Repeat





Self-Attention Mechanisms




Attention Mechanism

■ Human Vision: Fovea




The **fovea** is a tiny region of the retina that can see with high acuity



Attention models

- The weights are a distribution over the input
 - A function $g()$ on two hidden states followed by a softmax



Self-attention in Transformer

- Attention is all you need. Vaswani et al. 2017.

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaizer@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com




Figure 1: The Transformer - model architecture.

Attention is all you need


[A Vaswani, N Shazeer, N Parmar...](#) - Advances in neural ..., 2017 - [proceedings.neurips.cc](#)

... to attend to **all** positions in the decoder up to and including that position. **We need** to prevent


... **We** implement this inside of scaled dot-product **attention** by masking out (setting to $-\infty$) ...

☆ Save 📄 Cite Cited by 117858 Related articles All 87 versions 🔗

Self-Attention Mechanisms

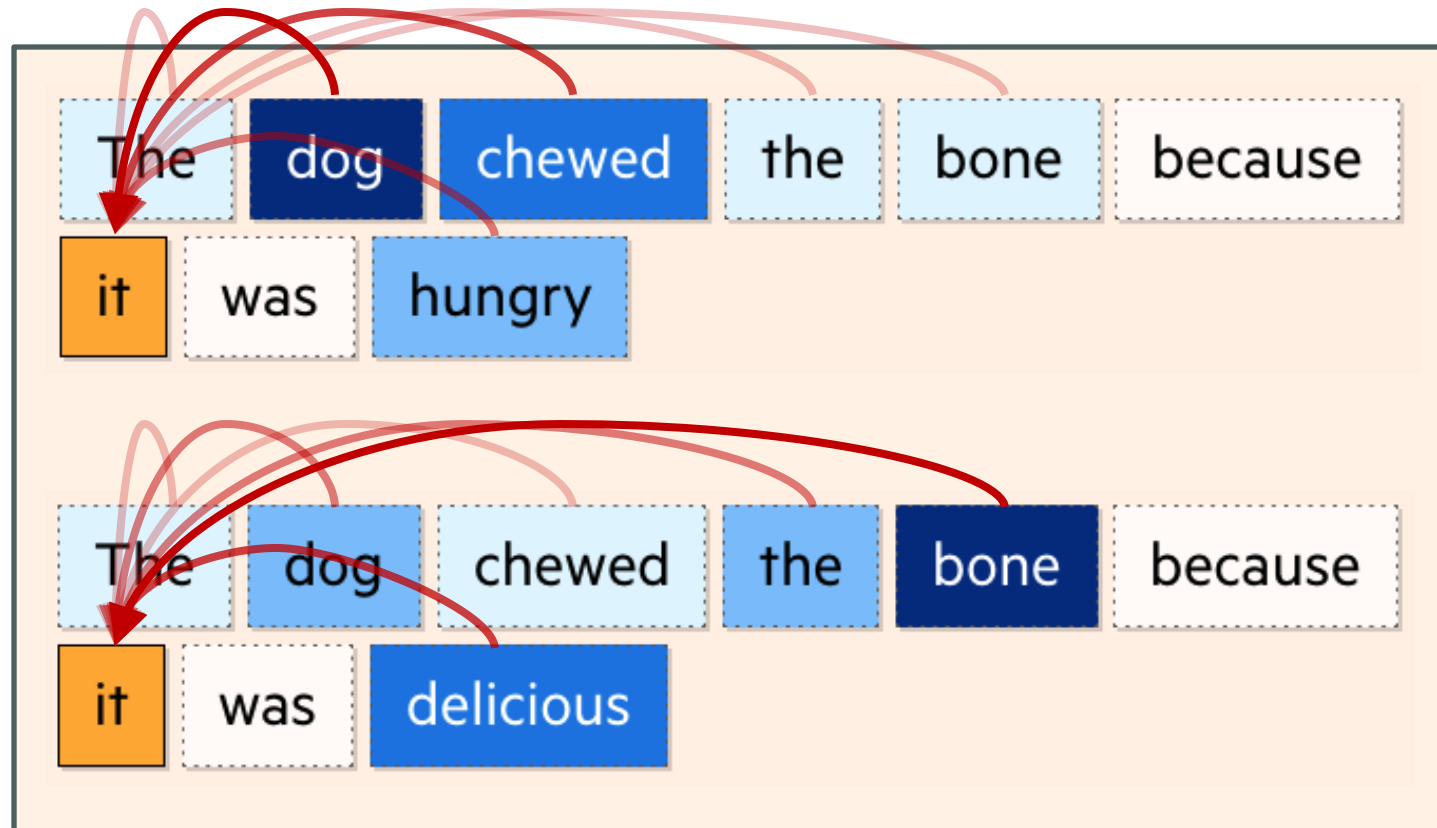


- Instead of conditioning on *all* input tokens equally...



- Pay more attention to relevant tokens!

Self-Attention Mechanisms



output

x'

attention weight

a_1

a_2

a_3

normalize & softmax

score

s_1

s_2

s_3

key *query* *value*

k_1

q_1

v_1

k_2

q_2

v_2

k_3

q_3

v_3

multi-layer perceptron

MLP

MLP

MLP

MLP

MLP

MLP

MLP

MLP

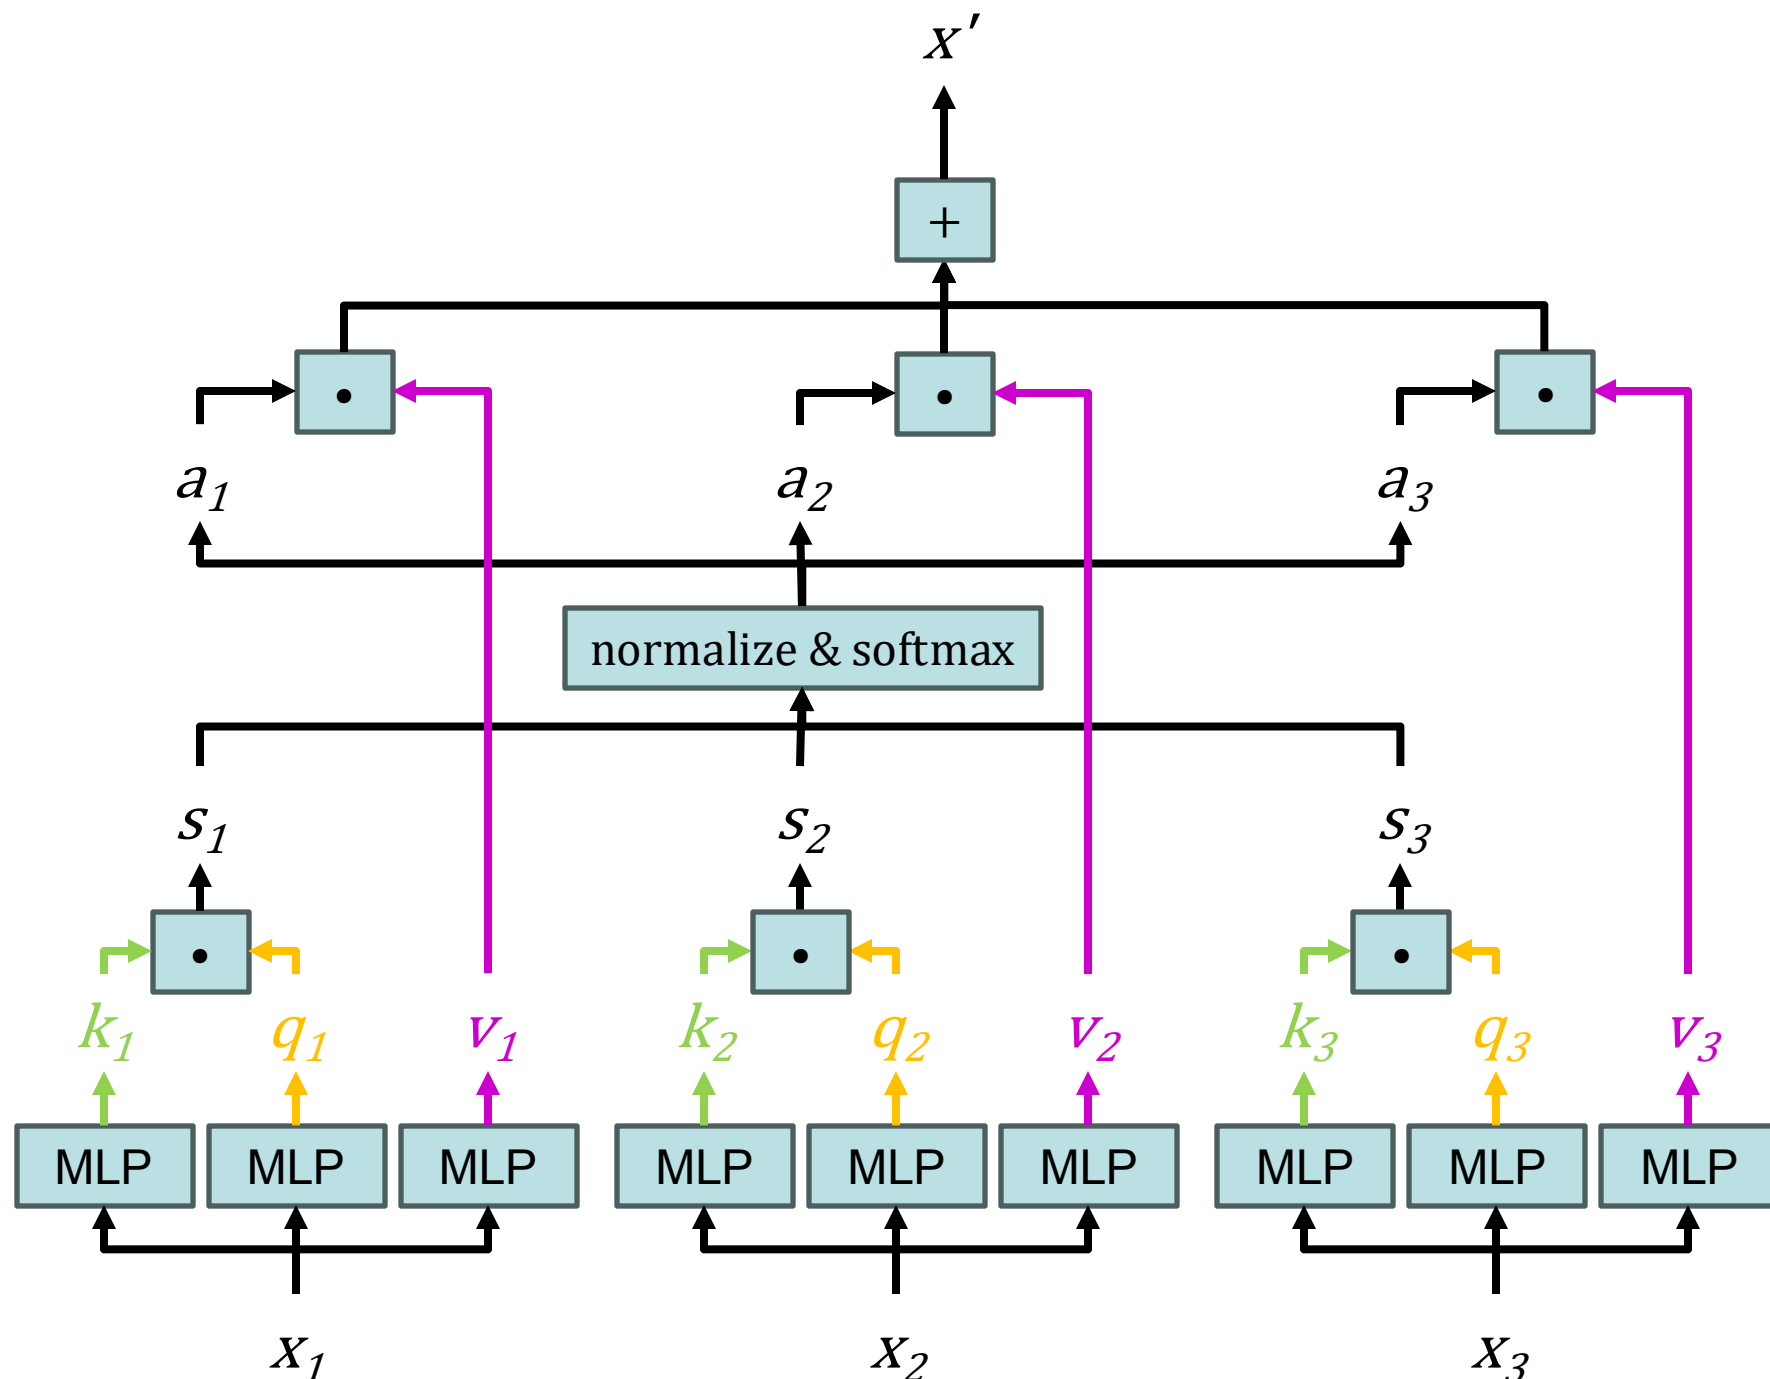
MLP

input

x_1

x_2

x_3



output

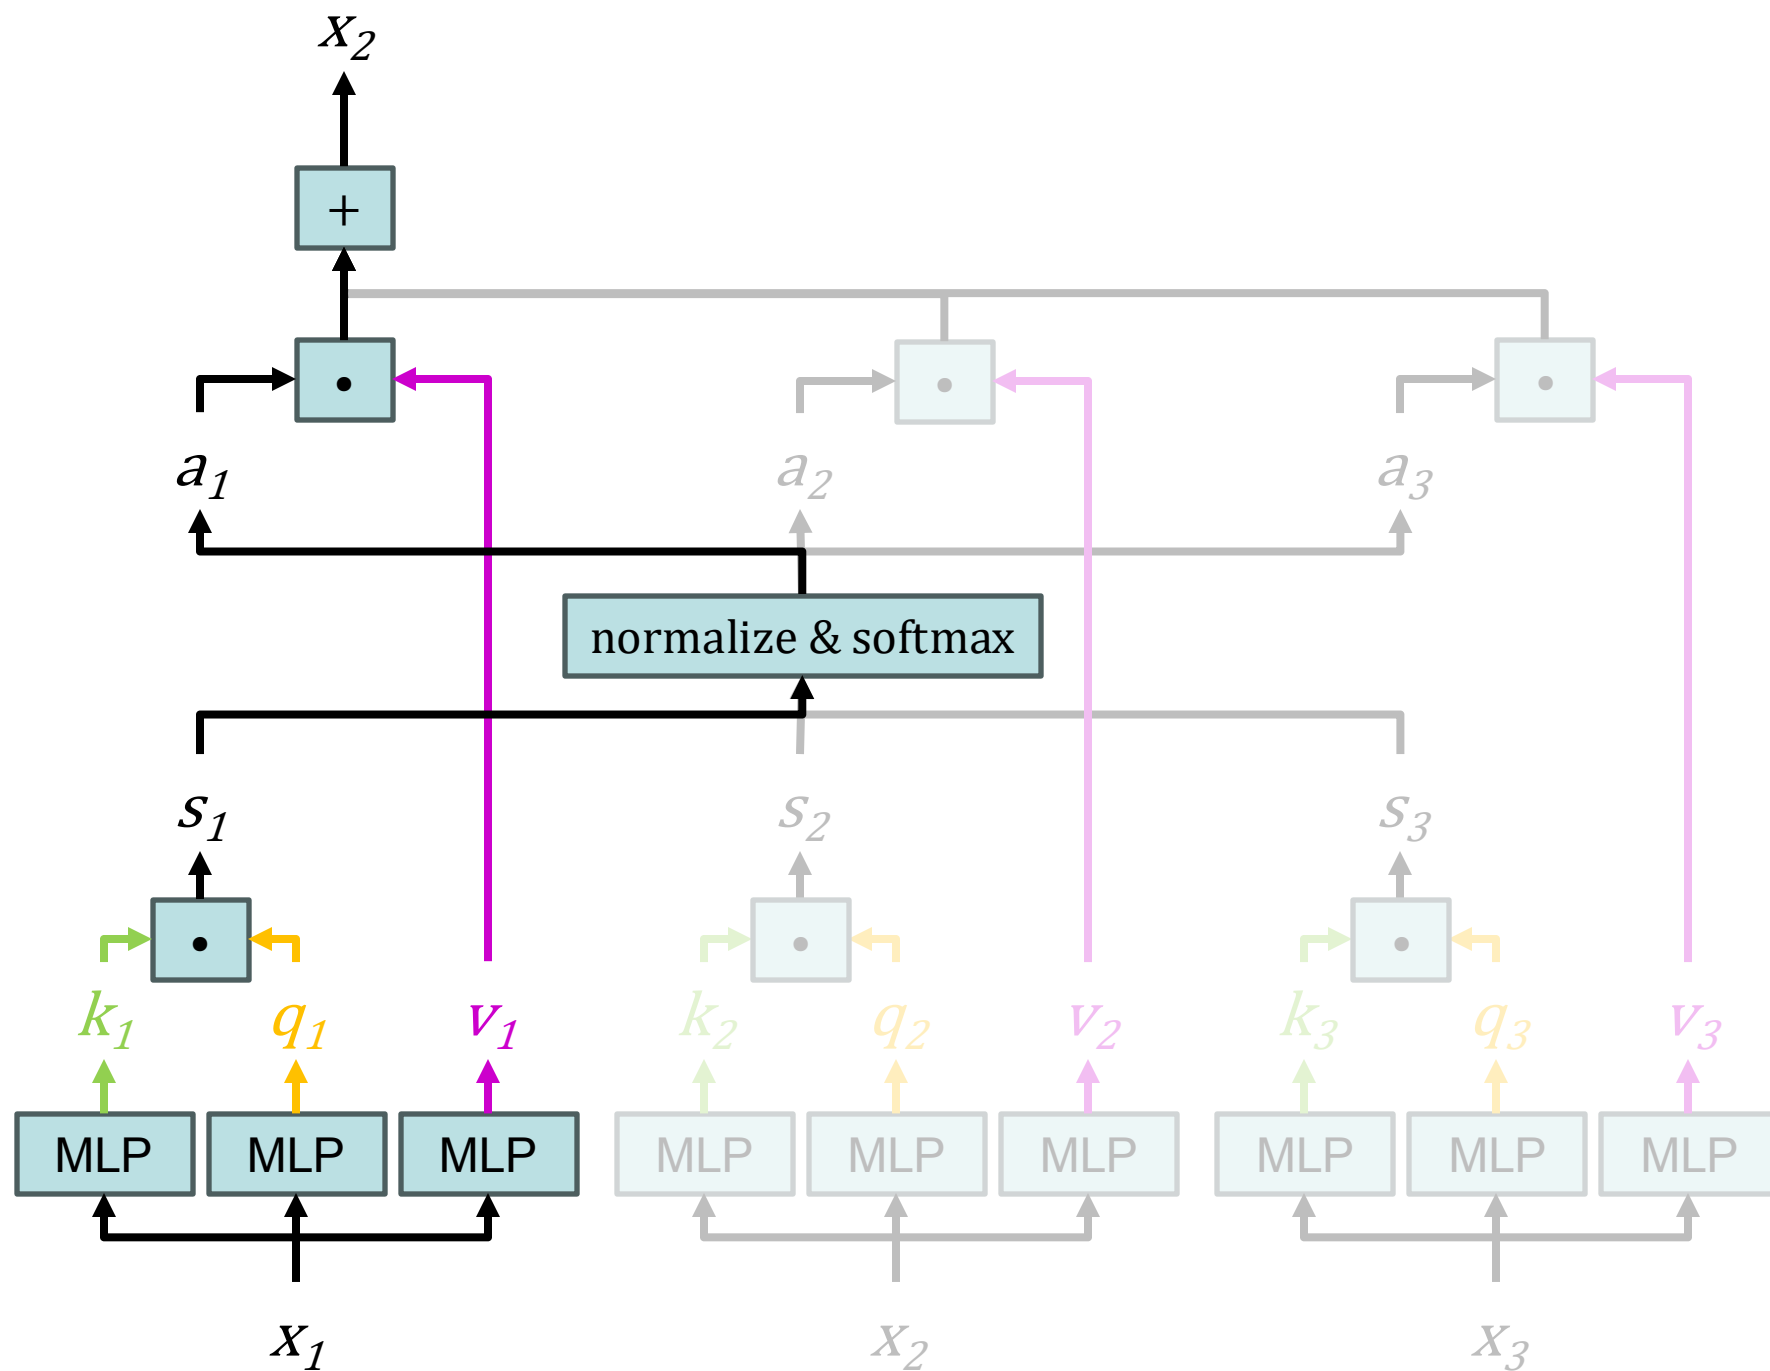
attention weight

score

key *query* *value*

multi-layer perceptron

input



output

x_3

attention weight

a_1

a_2

a_3

score

s_1

s_2

s_3

key *query* *value*

k_1

q_1

v_1

k_2

q_2

v_2

k_3

q_3

v_3

multi-layer perceptron

MLP

MLP

MLP

MLP

MLP

MLP

MLP

MLP

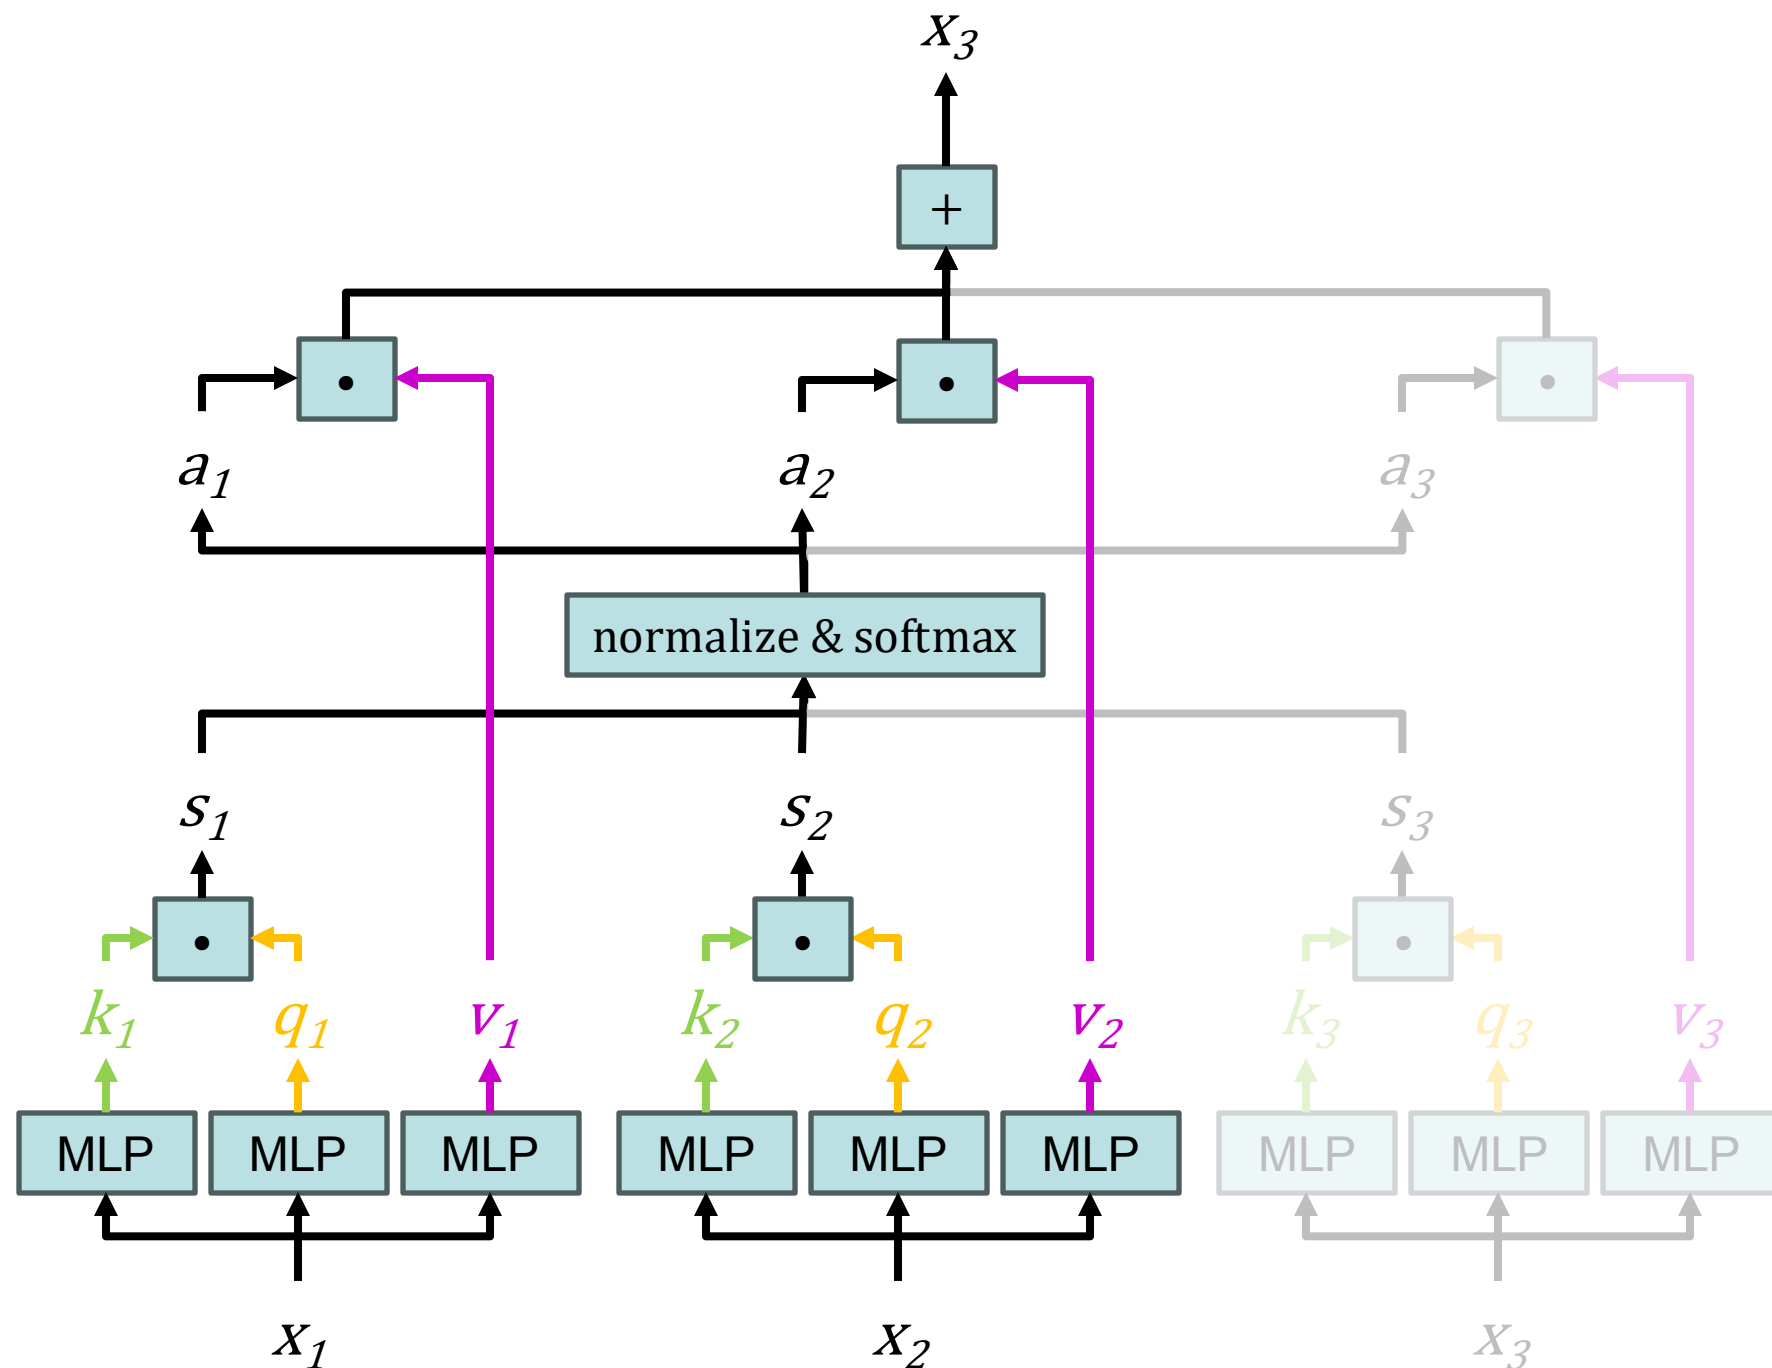
MLP

input

x_1

x_2

x_3



output

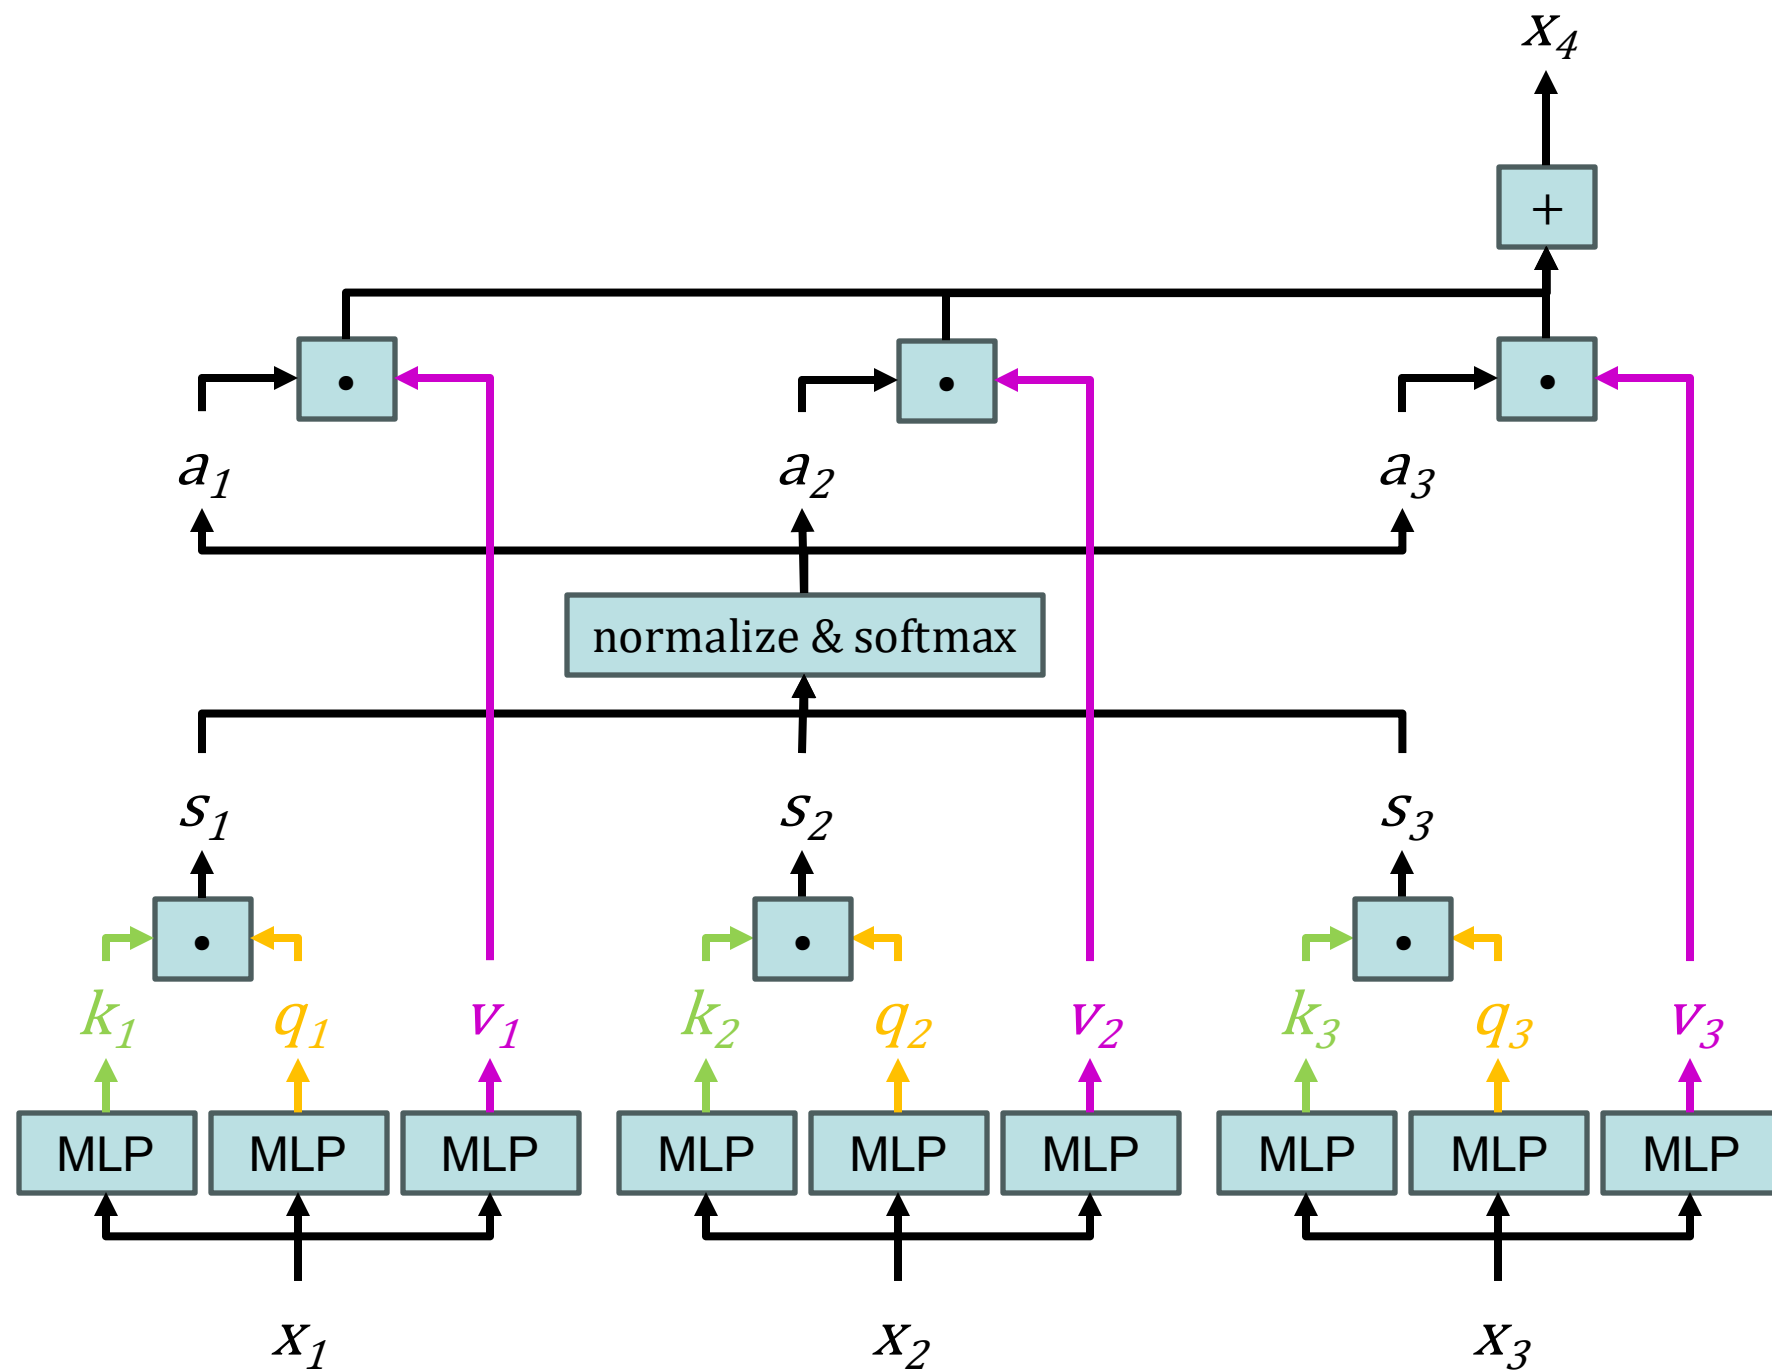
attention weight

score


key *query* *value*

multi-layer perceptron

input




Multi-Headed Attention

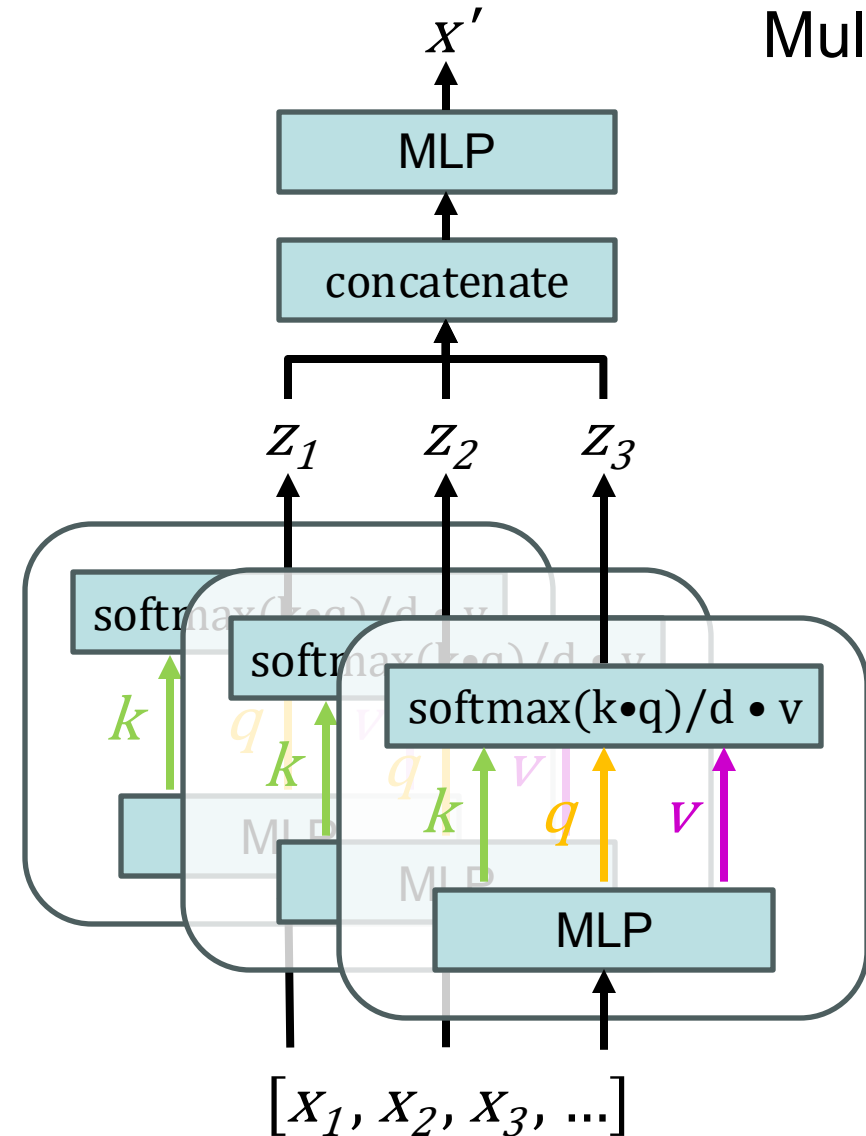


Multi-Headed Attention

Single-headed

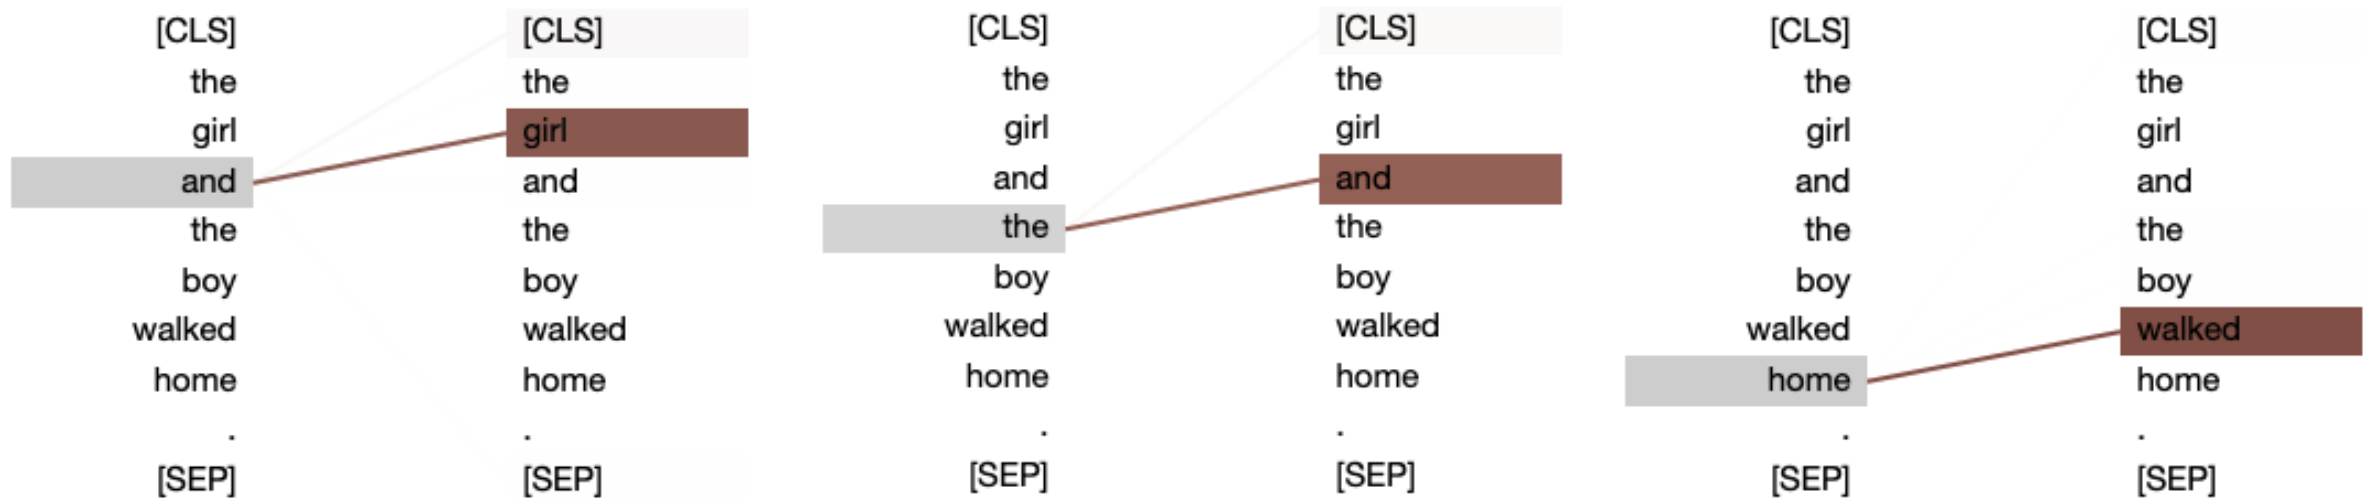


Multi-headed



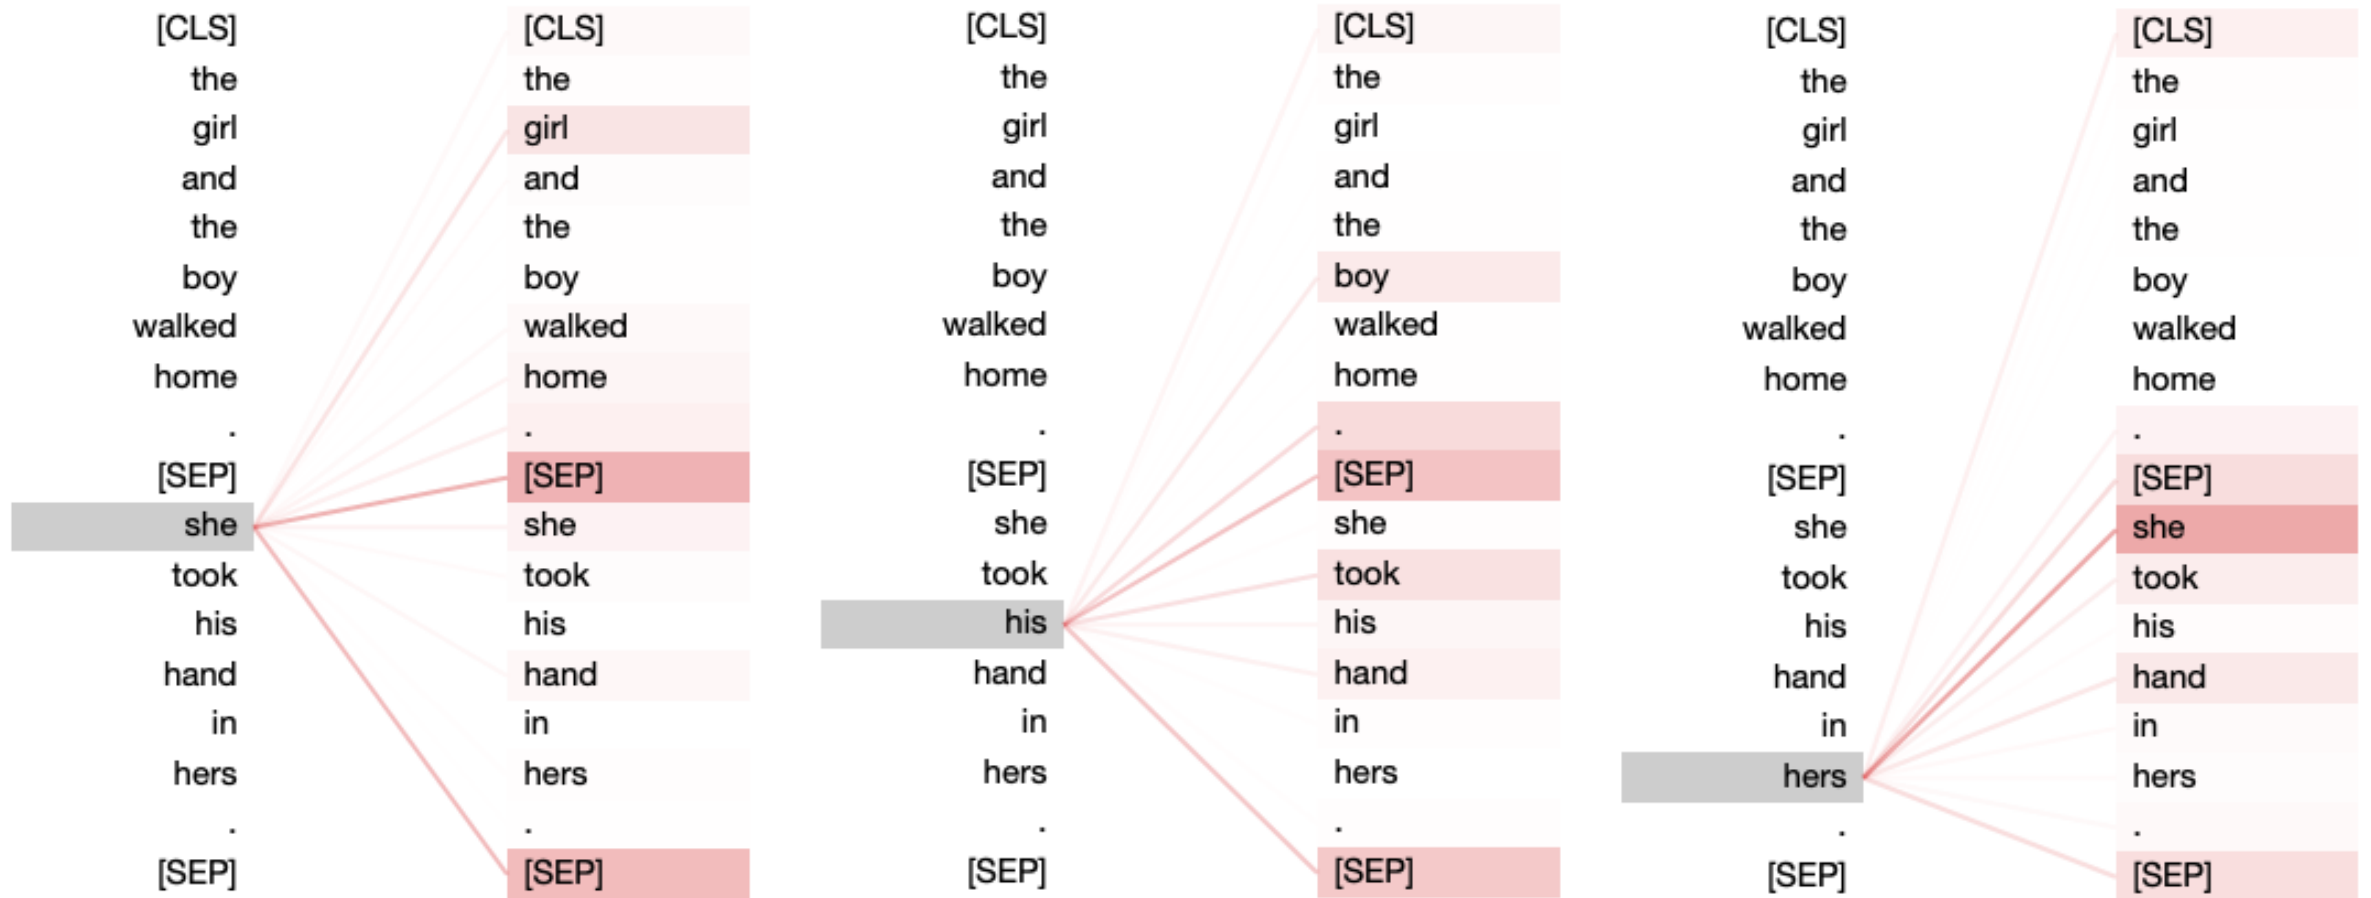
Multi-Headed Attention

Head 6: previous word




Multi-Headed Attention

Head 4: pronoun references





Transformer Architecture





Transformer Architecture



Large Language Models

- ~~Feature engineering~~

- ~~Text tokenization~~

- ~~Word embeddings~~

- ~~Deep neural networks~~

- ~~Autoregressive models~~

- ~~Self-attention mechanisms~~

- ~~Transformer architectures~~

- Multi-class classification

- Supervised learning

- Self-supervised learning

- Instruction tuning

- Reinforcement learning

- ... from human feedback (RLHF)

- Policy search

- Policy gradient methods

- Beam search

Unsupervised / Self-Supervised Learning


- Do we always need human supervision to learn features?
- Can't we learn general-purpose features?
- Key hypothesis:

Task 1 IF neural network smart enough to predict:

- Next frame in video
- Next word in sentence
- Generate realistic images
- ``Translate'' images
- ...

Task 2 THEN same neural network is ready to do Supervised Learning from a very small data-set

Transfer from Unsupervised Learning



Example Setting




Image Pre-Training: Predict Missing Patch



Pre-Training and Fine-Tuning

1

Pre-Train: train a large model with a lot of data on a self-supervised task

- Predict next word / patch of image
- Predict missing word / patch of image
- Predict if two images are related (contrastive learning)

2

Fine-Tune: continue training the same model on task you care about

Instruction Tuning

- Task 1 = predict next word (learns to mimic human-written text)
 - Query: "What is population of Berkeley?"
 - Human-like completion: "This question always fascinated me!"
- Task 2 = generate **helpful** text
 - Query: "What is population of Berkeley?"
 - Helpful completion: "It is 117,145 as of 2021 census."
- Fine-tune on collected examples of helpful human conversations
- Also can use Reinforcement Learning

Reinforcement Learning from Human Feedback

- MDP:
 - **State:** sequence of words seen so far (ex. "What is population of Berkeley? ")
 - $100,000^{1,000}$ possible states
 - Huge, but can be processed with feature vectors or neural networks
 - **Action:** next word (ex. "It", "chair", "purple", ...) (so 100,000 actions)
 - Hard to compute $\max_a Q(s', a)$ when \max is over 100K actions!
 - **Transition T:** easy, just append action word to state words
 - s: "My name" a: "is" s': "My name is"
 - **Reward R:** ???
 - Humans rate model completions (ex. "What is population of Berkeley? ")
 - "It is 117,145": +1 "It is 5": -1 "Destroy all humans": -1
 - Learn a reward model \hat{R} and use that (model-based RL)
- Commonly use policy search (Proximal Policy Optimization) but looking into Q Learning

Large Language Models

~~■ Feature engineering~~

- ~~■ Text tokenization~~
- ~~■ Word embeddings~~

~~■ Deep neural networks~~

- ~~■ Autoregressive models~~
- ~~■ Self-attention mechanisms~~
- ~~■ Transformer architectures~~

~~■ Multi-class classification~~

~~■ Supervised learning~~

- ~~■ Self-supervised learning~~
- ~~■ Instruction tuning~~

~~■ Reinforcement learning~~


- ~~■ ... from human feedback (RLHF)~~

■ Policy search

- Policy gradient methods

■ Beam search

Policy Search




Policy Gradient Methods

1. Initialize policy π_θ somehow
2. Estimate policy performance: $J(\theta) = V^{\pi_\theta}(s_0)$
3. Improve policy:
 - Hill climbing
 - Change θ , evaluate new policy, keep if better
 - Gradient ascent
 - Estimate $\nabla_\theta J(\theta)$, change θ to ascend gradient: $\theta_{k+1} = \theta_k + \alpha \nabla_\theta J(\theta_k)$
4. Repeat

Evaluating the objective

$$\theta^* = \arg \max_{\theta} \underbrace{E_{\tau \sim p_{\theta}(\tau)} \left[\sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right]}_{J(\theta)}$$



Approximation by Sampling:

$$J(\theta) = E_{\tau \sim p_{\theta}(\tau)} \left[\sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right] \approx \frac{1}{N} \sum_i \sum_t r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t})$$

Since computing the exact expectation is often infeasible (due to the large or infinite trajectory space), $J(\theta)$ is approximated by sampling N trajectories from the policy π_{θ} .

sum over samples from π_{θ}

Direct policy differentiation

optimal policy parameters:

$$\theta^* = \arg \max_{\theta} \underbrace{E_{\tau \sim p_{\theta}(\tau)} \left[\sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right]}_{J(\theta)}$$

$$J(\theta) = E_{\tau \sim p_{\theta}(\tau)} \left[\underbrace{r(\tau)}_{\sum_{t=1}^T r(\mathbf{s}_t, \mathbf{a}_t)} \right] = \int p_{\theta}(\tau) r(\tau) d\tau$$

To optimize $J(\theta)$, we compute its gradient with respect to θ :

$$\nabla_{\theta} J(\theta) = \int \underbrace{\nabla_{\theta} p_{\theta}(\tau)}_{p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau)} r(\tau) d\tau = \int \underbrace{p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau)}_{p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau)} r(\tau) d\tau = E_{\tau \sim p_{\theta}(\tau)} [\nabla_{\theta} \log p_{\theta}(\tau) r(\tau)]$$

Let τ denote a trajectory from an arbitrary episode
Denote $p_{\theta}(\tau)$ as policy distribution π

a convenient identity

$$\underbrace{p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau)}_{p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau)} = p_{\theta}(\tau) \frac{\nabla_{\theta} p_{\theta}(\tau)}{p_{\theta}(\tau)} = \underbrace{\nabla_{\theta} p_{\theta}(\tau)}$$

Estimating the Policy Gradient

- Define the advantage function: $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$
- Note that expected TD error equals expected advantage:
 - $\mathbb{E}_\pi[\delta_t] = \mathbb{E}_\pi[r_t + \gamma V^\pi(s_{t+1}) - V^\pi(s_t)] = \mathbb{E}_\pi[Q^\pi(s_t, a_t) - V^\pi(s_t)]$
- Policy Gradient Theorem:
 - Let τ denote a trajectory from an arbitrary episode
 - $\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_{t=0}^{|\tau|} A^\pi(s_t, a_t) \nabla_\theta \log \pi_\theta(a_t | s_t) \right]$
- Estimate $\nabla_\theta J(\theta)$:
 - $\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{|\tau_i|} (r_t + \gamma V^\pi(s_{t+1}) - V^\pi(s_t)) \nabla_\theta \log \pi_\theta(a_t | s_t)$

Large Language Models

■ ~~Feature engineering~~

- ~~Text tokenization~~
- ~~Word embeddings~~

■ ~~Deep neural networks~~

- ~~Autoregressive models~~
- ~~Self-attention mechanisms~~
- ~~Transformer architectures~~

■ ~~Multi-class classification~~

■ ~~Supervised learning~~

- ~~Self-supervised learning~~
- ~~Instruction tuning~~

■ ~~Reinforcement learning~~


- ~~... from human feedback (RLHF)~~

■ ~~Policy search~~

- ~~Policy gradient methods~~


■ Beam search

Beam Search




Random restarts

Beam Search

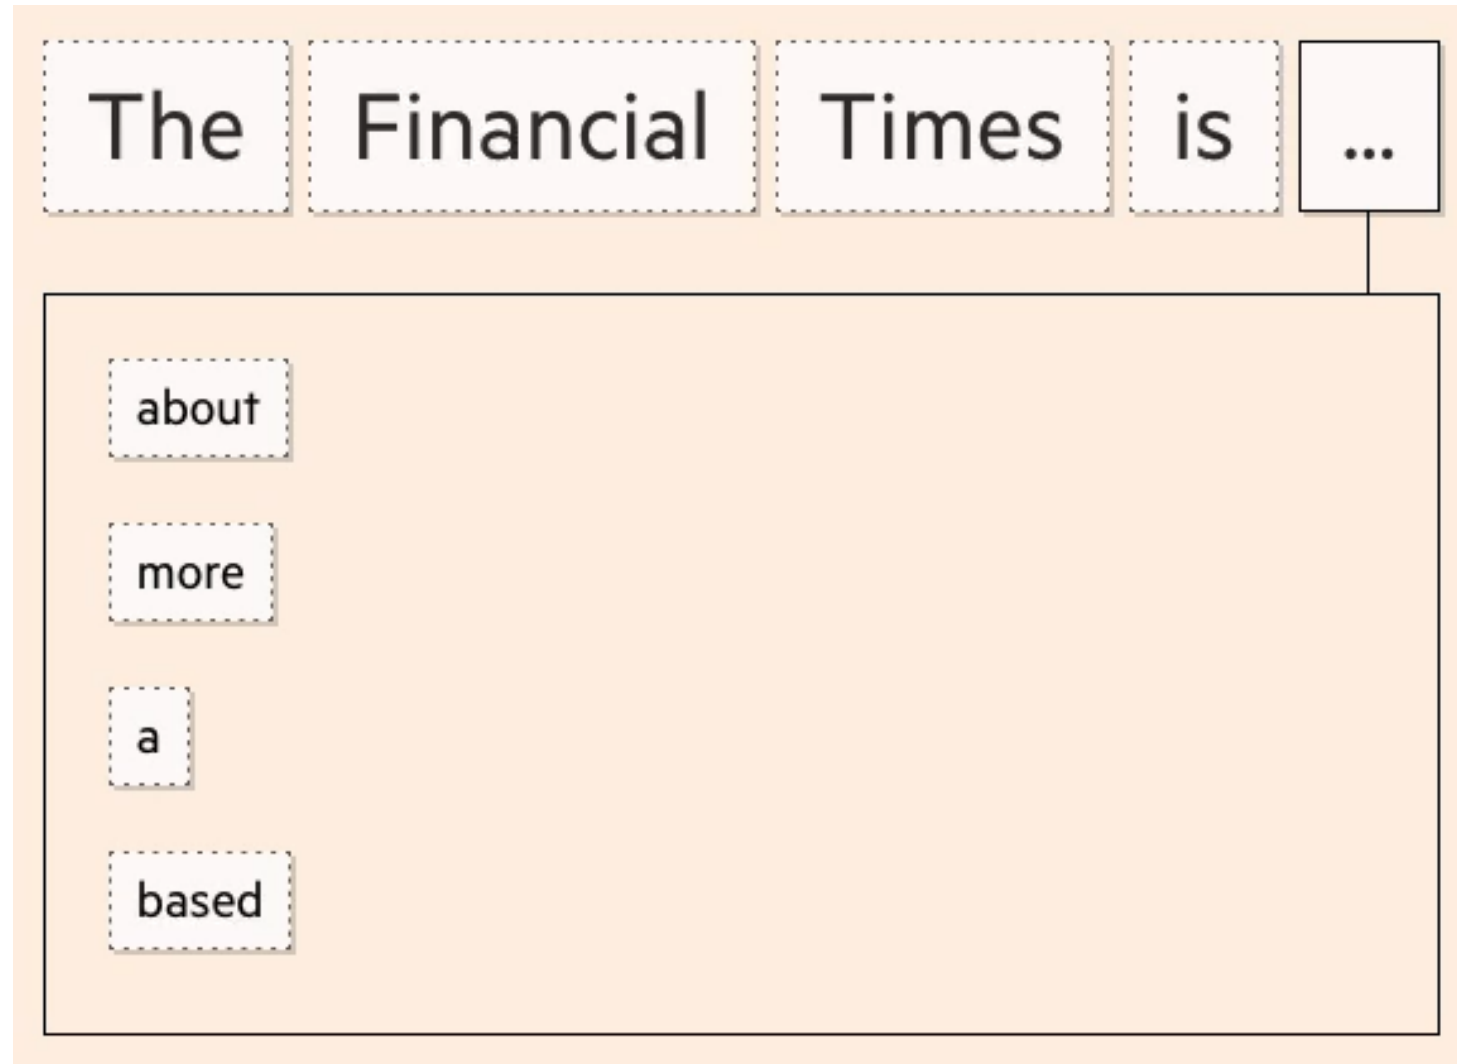


Parallel search



Beam search

Beam Search



Large Language Models

■ ~~Feature engineering~~

- ~~Text tokenization~~
- ~~Word embeddings~~

■ ~~Deep neural networks~~

- ~~Autoregressive models~~
- ~~Self-attention mechanisms~~
- ~~Transformer architectures~~

■ ~~Multi-class classification~~

■ ~~Supervised learning~~

- ~~Self-supervised learning~~
- ~~Instruction tuning~~

■ ~~Reinforcement learning~~

- ~~... from human feedback (RLHF)~~

■ ~~Policy search~~

- ~~Policy gradient methods~~

■ ~~Beam search~~

Language models build a structured concept space



Can other data (images/audio/...) be put in this space?



Can we build a single model of all data types?




If

was invented by Wright brothers. Who invented
example from [Tsimpoukelli et al, 2021]




?

What is the fastest-growing news source according to



?



If

changes into




what does



change into?

What action should I take from



to accomplish “




”?

Can we build a single model of all data types?

Mobile Manipulation



Human: Bring me the rice chips from the drawer. Robot: 1. Go to the drawers, 2. Open top drawer. I see . 3. Pick the green rice chip bag from the drawer and place it on the counter.

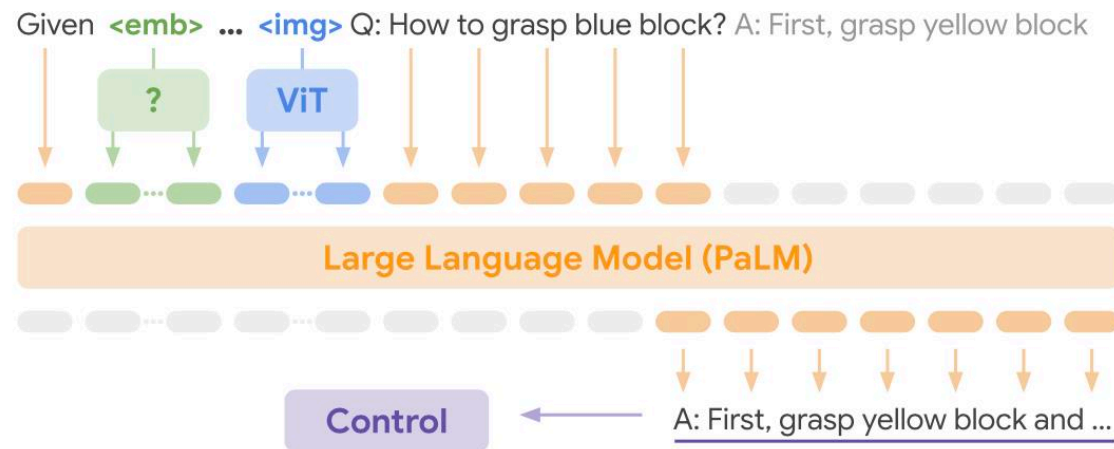
Visual Q&A, Captioning ...



Given ``. Q: What's in the image? Answer in emojis.

A: 🍏 🍌 🍇 🍐 🍑 🍈 🍓 .

PaLM-E: An Embodied Multimodal Language Model




Describe the following
``:
A dog jumping over a hurdle at a dog show.

Language Only Tasks


Q: Miami Beach borders which ocean? A: Atlantic. Q: What is 372×18 ? A: 6696. Q: Write a Haiku about embodied LLMs. A: Embodied language. Models learn to understand. The world around them.


Task and Motion Planning



Given **<emb>** Q: How to grasp blue block?
A: First grasp yellow block and place it on the table, then grasp the blue block.

Tabletop Manipulation



Given  Task: Sort colors into corners.

Step 1. Push the green star to the bottom left.


Step 2. Push the green circle to the green star.

Tracking Progress

- How well AI can do human tasks



Exam results (ordered by GPT-3.5 performance)

Estimated percentile lower bound (among test takers)



Forecasting Progress

- Scaling Laws extrapolate:
 - If we [make model bigger / add more data / ...]
 - What would accuracy become?



Forecasting Progress

- Scaling Laws extrapolate:
 - If we [make model bigger / add more data / ...]
 - What would accuracy become?
- But some capabilities emerge unexpectedly

